

AEAF

Architecture Definition Document

Parkki Nordic agentic support line — BDAT+I

Date: 03.06.2026


Contents

Architecture overview	4
Business	6
Data	6
Application	6
Technology	6
Intelligence & Agency (+I)	7
Key decisions	7
Compliance & gaps	7
Sign-off	7

Architecture Definition Document — Parkki Nordic agentic support line

Deliverable (the headline board document). The single composition of the whole target architecture across BDAT+I. It assembles the domain artifacts with a synthesis per layer; it does not re-author them — every claim traces to the artifact that holds the detail.

Architecture overview



One picture of the whole: the four classical domains carry the service; the +I domain carries the agents, models, guardrails, evals, and autonomy; human oversight and accountability wrap it.

Business

The capability model decomposed to the decision level, with performer mix and autonomy per leaf (→ [business-capability-catalog](#)); the responsibility/accountability split and decision bands (→ [human-agent-raci](#)); the actors and roles (→ [organization-actor-catalog](#) , [role-catalog](#)); and the frontier and its trajectory (→ [capability-automation-frontier-map](#)). Refunds split at the decision level; safety and account/GDPR fixed human.

Data

Data entities with their classification and the [feedsCorpus](#) link (→ [data-entity-catalog](#)), and the four governed knowledge corpora with freshness, classification, and owner (→ [knowledge-corpus-catalog](#)). The load-bearing decision: **no customer PII enters a corpus** — policy is embedded, customer facts are read live.

Application

The agent runtime, voice front-end, and core platform (→ [application-portfolio-catalog](#)), and the non-determinism design at the process level — bounded outputs, fallback, handoff, and "deny and escalate, never open on uncertainty" for the gate (→ [process-function-catalog](#)).

Technology

The AI platform with its enforcement points named (→ [technology-portfolio-catalog](#)): the model gateway (PII egress, [6-3](#)), the orchestration runtime (pre-action guardrails), the eval pipeline, and the WORM audit store. The reasoning models as a right-sized portfolio (→ [model-portfolio](#)).

Intelligence & Agency (+I)

The five agents with their bounded purposes (→ [agent-catalog](#)); the guardrails compiled from principles, with enforcement points and fail actions (→ [guardrail-policy-catalog](#)); the eval suites with thresholds fixed before the run (→ [eval-suite-specification](#)); the autonomy assignment against the canonical scale (→ [autonomy-level-classification](#)); and the live standing view of who is answerable now (→ [trust-accountability-matrix](#)).

Key decisions

- [ADR-001](#) — Gate-Ops goes live at **L2**, **not L3** (physical action, no live evidence yet).
- [ADR-002](#) — EU AI Act classification: **limited-risk** (counsel-confirmed).

Compliance & gaps

The EU obligations mapped to the controls that meet them (→ [regulatory-obligations-register](#)): GDPR (incl. Art 22 via the autonomy model), AI Act Art 50 transparency, ePrivacy consent. One open item — accessibility — tracked as partial. Baseline-to-target gaps and the sequence to close them are the Roadmap (→ [architecture-roadmap-migration-plan](#)).

Sign-off

Architecture board, 2026-04-02. This ADD is the agreed target architecture; the design-time gate (per agent) authorises operation.