

AEAF

Architecture Vision

Parkki Nordic agentic support line

Date: 09.06.2026

Contents

Problem & drivers	3
Stakeholders & concerns	3
Human-agent ambition (+I)	3
Consequence / risk envelope (+I)	4
Target value	4
Headline intent	4
Guiding principles	4
Target frontier	4
Sign-off	5

Architecture Vision — Parkki Nordic agentic support line

Deliverable (Phase 1). Secures agreement on where the support line is going and how far autonomy may go. Assembled from the intent records, the principles, and the target frontier; signed by the architecture board.

Problem & drivers

Support volume tripled as the garage estate grew. The aim is to resolve routine, reversible contacts (stuck at a paid gate, "why was I charged", a camera misread) in minutes, while keeping sensitive, contested, and irreversible work with people. The board's condition: this runs real money and a physical gate, so it must be **governable**, not a demo.

Stakeholders & concerns

Drivers stuck at exit; the human support team (whose work changes); the Billing and Data functions; the DPO (personal/payment data); the board (risk and lawfulness). The people whose roles change are in scope of the Vision, not an afterthought.

Human-agent ambition (+I)

A blended workforce in front of the human support team: **agents resolve the routine, reversible contacts** (stuck at a paid gate, a charge query, a camera misread), and **humans own the sensitive and the irreversible**. How that work decomposes into specific agents, and at what autonomy, is the Phase 2 allocation and the Phase 6 agent design — not something the Vision fixes here.

Consequence / risk envelope (+I)

The board caps the support line at **autonomy L3** (no out-of-the-loop action — every case touches money or a gate). Three classes of work are **fixed human**: safety/emergency, account/GDPR, and large or contested money. These are decisions, not gaps.

Target value

Routine contacts resolved in minutes; the human team freed for the cases that need judgment; quality, safety, and continuity provable to the board and a regulator. The latency, resolution rate, and cost this yields are a technology outcome, designed and measured downstream — not numbers committed at the Vision.

Headline intent

The agent work is declared as **intent** — outcome + bounds, assured by evals, not a fully-specified script (→ [requirements-intent-catalog](#)). The specific intent records and their thresholds are elaborated as the capabilities are allocated and the eval suites designed.

Guiding principles

Principles [P-1...P-10](#) (→ [principles-catalog](#)) — five runtime-enforced (PII boundary, safety-to-human, verified-facts-only, bounded money, paid-session-for-gate), plus the GDPR/AI Act/ePrivacy principles. Each runtime principle compiles to a guardrail.

Target frontier

The intended direction: **agents take the routine and reversible work; the boundary widens only as assurance accrues; and three classes stay human — safety, account/GDPR, and large or contested money**. The per-capability boundary and its trajectory are mapped in Phase 2 (→ [capability-automation-frontier-map](#)), not pre-drawn here.

Sign-off

Architecture board (chair: Head of Customer Support), 2026-02-20. Proceed to Phase 2.