

AEAF

Architecture Vision — Parkki Nordic agentic support line

Phase 1 deliverable

Date: 03.06.2026

Contents

Problem & drivers	3
Stakeholders & concerns	3
Human-agent ambition (+I)	3
Consequence / risk envelope (+I)	3
Target value	4
Headline intent	4
Guiding principles	4
Target frontier	4
Sign-off	4

Architecture Vision — Parkki Nordic agentic support line

Deliverable (Phase 1). Secures agreement on where the support line is going and how far autonomy may go. Assembled from the intent records, the principles, and the target frontier; signed by the architecture board.

Problem & drivers

Support volume tripled as the garage estate grew. The aim is to resolve routine, reversible contacts (stuck at a paid gate, "why was I charged", a camera misread) in minutes, while keeping sensitive, contested, and irreversible work with people. The board's condition: this runs real money and a physical gate, so it must be **governable**, not a demo.

Stakeholders & concerns

Drivers stuck at exit; the human support team (whose work changes); the Billing and Data functions; the DPO (personal/payment data); the board (risk and lawfulness). The people whose roles change are in scope of the Vision, not an afterthought.

Human-agent ambition (+I)

A five-agent blended workforce — Triage, Gate-Ops, Billing-Dispute, ANPR-Correction, Escalation-Router — in front of the human team. Agents resolve the routine; humans own the sensitive and the irreversible.

Consequence / risk envelope (+I)

The board caps the support line at **autonomy L3** (no out-of-the-loop action — every case touches money or a gate). Three classes of work are **fixed human**: safety/emergency, account/GDPR, and large or contested money. These are decisions, not gaps.

Target value

Routine contacts resolved in under ~90 seconds; the human team freed for the cases that need judgment; quality, safety, and continuity provable to the board and a regulator.

Headline intent

The five agent intents `INT-001...INT-005` (→ `requirements-intent-catalog`) — each an outcome + bounds, assured by an eval suite, not a fully-specified script.

Guiding principles

Principles `P-1...P-10` (→ `principles-catalog`) — five runtime-enforced (PII boundary, safety-to-human, verified-facts-only, bounded money, paid-session-for-gate), plus the GDPR/AI Act/ePrivacy principles. Each runtime principle compiles to a guardrail.

Target frontier

The intended human/agent boundary and its trajectory (→ `capability-automation-frontier-map`): agents on triage/explain/gate/ANPR with a path to widen on evidence; humans fixed on safety, account/GDPR, and large/contested money.

Sign-off

Architecture board (chair: Head of Customer Support), 2026-02-20. Proceed to Phase 2.