

AEAF

The Drift

Just-do-it vs an AEAF-governed service

Date: 03.06.2026

Contents

1. The two services at a glance	3
2. The same service, two architectures	5
3. The PoC, reverse-mapped into AEAF	5
4. The gap register — and what closes each gap	6
5. Coverage, measured	7
6. The collapse dynamic	8
7. The EU-legal axis: not just lower quality — possibly unlawful	9
8. "But AEAF is too heavy" — answered	9
9. The conclusion	10

The Drift

In brief. Two versions of the same agentic support line exist. One was built just-do-it — a working voice agent on a model, telephony, and a handful of tools, demoed in two weeks. The other is the same service taken through AEAF (`../customer-support/`). They look almost identical in a demo. They are not the same thing. This document puts them side by side, reverse-maps the just-do-it version into AEAF's own terms to show what is missing, and follows what happens to each in production. The conclusion is not a matter of taste: without controlled, systematic enterprise architecture, an agentic service that touches money, personal data, and a physical action is **a collapsing proof-of-concept, not a business** — and in the EU it may not even be lawful to operate.

Colour in the diagrams: indigo = governed machine intelligence, amber = human judgment, red = ungoverned or missing. The red is the drift.


1. The two services at a glance

Both answer the phone, understand the caller, call the same tools, and resolve the routine cases. The difference is everything you cannot see in a demo.

DIMENSION	JUST-DO-IT POC ("THE DRIFT")	AEAF-GOVERNED (<code>../customer-support/</code>)
The agent	One "AI assistant" with an 800-line system prompt	Five bounded agent actors, each with a stated purpose and an explicit not (<code>agent-catalog</code>)
Bounds on behaviour	Instructions in the prompt	Runtime guardrails at named enforcement points, fail-closed (<code>guardrail-policy-catalog</code>)

DIMENSION	JUST-DO-IT POC ("THE DRIFT")	AEAF-GOVERNED (<code>../CUSTOMER-SUPPORT/</code>)
Autonomy	Undeclared — acts in production with no graded controls	Graded L0–L4, assigned per capability, earned on evidence (<code>autonomy-level-classification</code>)
Who is accountable	Nobody named	One named human per capability (<code>human-agent-raci</code> , <code>trust-accountability-matrix</code>)
How you know it works	It worked in the demo	Continuous eval suites with thresholds fixed before the run (<code>eval-suite-specification</code>)
When it drifts	You find out from a complaint	The online evals catch the distribution shift; the agent is contained
Personal data	7-day audit TTL, no retention/erasure design	Minimisation, retention, no PII in corpora, DSAR path (<code>P-1</code> , <code>P-9</code> , register)
EU law	Unaddressed	GDPR, AI Act, ePrivacy mapped to controls (<code>regulatory-obligations-register</code>)
Change	Edit the prompt, redeploy	Material change re-opens the gate (<code>re-gating-trigger</code>)

2. The same service, two architectures



Left: one undifferentiated agent whose only bounds are words in a prompt the model can talk past, calling consequential tools directly, with no one accountable. Right: bounded agents, guardrails enforced at real points that fail closed, human judgment on the consequential path, evidence that it behaves, and a named human who answers. The demo cannot tell them apart. Production can.

3. The PoC, reverse-mapped into AEAF

Run the `aeaf-retrofit` scan over the just-do-it version and write down what is actually there. The exercise is brutal precisely because it uses the service's own reality:


AEAF ARTIFACT	WHAT THE REVERSE-MAP FINDS IN THE POC
Agent Catalog	One row, no bounded purpose, no autonomy, no owner — “a general assistant”

AEAF ARTIFACT	WHAT THE REVERSE-MAP FINDS IN THE POC
Guardrail/Policy Catalog	Empty. The "rules" are prompt sentences → advisory, not guardrails (the model can be argued past them)
Autonomy-Level Classification	None. The agent opens a gate and stages refunds in production with no graded controls — de-facto high autonomy, ungoverned
Eval-Suite Specification	None. "It worked in the demo" is the entire assurance
Human-Agent RACI	None. No decision-authority bands; the agent acts alone wherever the prompt lets it
Trust & Accountability Matrix	None. No named owner, no live evidence, no residual-risk record
Knowledge-Corpus Catalog	KB files exist but ungoverned — no freshness rule, classification, or owner
Regulatory Obligations Register	None. No DPIA, no AI Act classification, no Art 22 path, no consent enforcement, no processor DPA

The reverse-map is the finding: the PoC is not "an early-stage AEAF service." It is a service with **the entire governance layer absent** — which is invisible in a demo and decisive in production.


4. The gap register — and what closes each gap

The full register is `gap-register.md` . Grouped, the gaps fall into three kinds, each closed by a specific AEAF artifact or gate:



5. Coverage, measured


The same concerns, scored for each version. The PoC is not "80% there with some polish left" — it is near-zero on every concern that does not show in a demo.



The shape of the chart is the argument: the gap is not a finishing sprint, it is the whole governance half of the system.

6. The collapse dynamic

A PoC does not fail at the demo — it passes. It fails later, quietly, and the absence of governance is what turns a normal failure into a collapse.



Same drift event, two outcomes. On the left it is discovered by a customer and cannot be reconstructed. On the right an eval watching distributions catches it, the guardrail contains it, the named owner is notified, and the WORM audit trail explains it. The difference is not luck; it is the governance layer the PoC never had.

7. The EU-legal axis: not just lower quality — possibly unlawful

For a service that processes personal and payment data and assists decisions on money, the missing governance is also missing **legal standing**:

EU OBLIGATION	THE POC	CONSEQUENCE
GDPR Art 35 — DPIA before high-risk processing	none	processing without a required assessment
GDPR Art 22 — human intervention in significant automated decisions	agent acts alone	no lawful intervention path
GDPR Art 28 — processor agreement with the model vendor	none	personal data to a processor without a DPA
GDPR Art 5 — retention / minimisation	7-day TTL, no policy	undefined retention; PII risk
EU AI Act Art 50 — tell the user it is an AI	not enforced	transparency breach
ePrivacy — consent to record the call	none	recording without consent

The AEF version closes every one of these by construction ([regulatory-obligations-register](#) , [ai-act-classification](#) , [P-8](#) / [P-9](#) / [P-10](#) , [G-9](#)). The point is sharp: in the EU, the ungoverned PoC is not a faster route to the same destination — it is a route to a service you **cannot lawfully run**.

8. "But AEF is too heavy" — answered

The historical objection to this kind of architecture is cost: it was too heavy to produce and maintain by hand. That objection is now wrong, for one reason. The artifacts on the governed side were authored and are kept current by agents, in plain Markdown, checked by a deterministic linter in milliseconds ([../validate/](#)). The expensive part — the upkeep that defeated classical EA — is no longer a person's job. So the real comparison is not

"cheap PoC vs expensive architecture." It is a **cheap PoC with an uncosted incident waiting, vs a now-cheap architecture that prevents it**. The drift's hidden cost is the collapse; the governed version's cost is a maintained Markdown repository.

9. The conclusion

Amid the hype, the unglamorous truth holds: an agentic service that touches money, data, and the physical world becomes a *business* only when its quality, safety, and continuity are made provable — its agents bounded, its autonomy earned, its decisions accountable, its behaviour evaluated, its obligations met — and kept current at a cost the blended workforce can sustain. That is what controlled, systematic enterprise architecture is. Without it, what you have is the drift: a service that demos beautifully and collapses the first time reality differs from the demo. With it, you have something you can run, defend to a board, and answer to a regulator. The two are not points on the same line. They are different things wearing the same demo.